

pytorch-basic-neural-networks

December 17, 2025

1 Lab 4: Pytorch - basic neural networks

1.0.1 Full Name: Nguyen Anh Thư

1.0.2 Student's ID: 23127266

2 PyTorch Feed Forward Neural Network

2.1 STEP 1 — Environment Setup

Import required libraries: `torch`, `numpy`

```
[34]: # Uncomment if required to install dependencies for this lab
```

```
#!pip -q install torch numpy
```

```
[35]: import torch
import numpy as np
import torch.nn as nn

print(f"PyTorch version: {torch.__version__}")
print(f"NumPy version: {np.__version__}")
```

PyTorch version: 2.7.1+cpu

NumPy version: 2.2.3

2.2 STEP 2 — Working with Tensors

Perform conversions between NumPy arrays and PyTorch Tensors, demonstrating shared memory mechanism.

```
[36]: # Create NumPy array
np_array = np.array([[1, 2, 3], [4, 5, 6]], dtype=np.float32)
print("NumPy array:")
print(np_array)

# Convert NumPy → PyTorch Tensor
tensor_from_np = torch.from_numpy(np_array)
print("\nPyTorch Tensor from NumPy:")
```

```

print(tensor_from_np)

# Demonstrate shared memory: changing NumPy array will affect Tensor
np_array[0, 0] = 999
print("\nAfter modifying NumPy array:")
print("NumPy array:", np_array)
print("PyTorch Tensor (shared memory):", tensor_from_np)

# Convert PyTorch Tensor → NumPy
np_array_from_tensor = tensor_from_np.numpy()
print("\nNumPy array from Tensor:")
print(np_array_from_tensor)

# Perform addition
result = np_array + np_array_from_tensor
print("\nAddition result:")
print(result)

```

NumPy array:
[[1. 2. 3.]
 [4. 5. 6.]]

PyTorch Tensor from NumPy:
tensor([[1., 2., 3.],
 [4., 5., 6.]])

After modifying NumPy array:
NumPy array: [[999. 2. 3.]
 [4. 5. 6.]]
PyTorch Tensor (shared memory): tensor([[999., 2., 3.],
 [4., 5., 6.]])

NumPy array from Tensor:
[[999. 2. 3.]
 [4. 5. 6.]]

Addition result:
[[1998. 4. 6.]
 [8. 10. 12.]])

2.3 STEP 3 — Define Activation Functions

Class ActivationFunction contains activation functions and their derivatives.

[37]: `class ActivationFunction:`
 `"""Utility class for activation functions and their derivatives.`

This class provides static methods for sigmoid activation function

```

and its derivative, using PyTorch operations only.

"""

@staticmethod
def sigmoid(s):
    """Compute sigmoid activation function.

Args:
    s: Input tensor

Returns:
    Tensor: Sigmoid activation of input, output range [0, 1]
"""
    return 1 / (1 + torch.exp(-s))

@staticmethod
def sigmoid_derivative(s):
    """Compute derivative of sigmoid function.

Args:
    s: Input tensor (typically sigmoid output)

Returns:
    Tensor: Derivative of sigmoid at input point
"""
    return s * (1 - s)

```

```
[38]: # Test activation functions
test_input = torch.tensor([0.0, 1.0, -1.0], dtype=torch.float32)
print("Input:", test_input)
print("Sigmoid:", ActivationFunction.sigmoid(test_input))
print("Sigmoid derivative:", ActivationFunction.
      sigmoid_derivative(ActivationFunction.sigmoid(test_input)))
```

```
Input: tensor([ 0.,  1., -1.])
Sigmoid: tensor([0.5000, 0.7311, 0.2689])
Sigmoid derivative: tensor([0.2500, 0.1966, 0.1966])
```

2.4 STEP 4-8 — Define Feed Forward Neural Network

Class `FFNeuralNetwork` inherits from `nn.Module`, containing:

- Initialization with weights and intermediate variables
- Forward propagation
- Backward propagation
- Training function
- Save/Load weights

```
[39]: class FFNeuralNetwork(nn.Module):
    """Feed Forward Neural Network implementation from scratch.
```

A simple 3-layer neural network (input -> hidden -> output) with

sigmoid activation function and manual backpropagation implementation.

Attributes:

inputSize: Number of input features (default: 3)
hiddenSize: Number of neurons in hidden layer (default: 4)
outputSize: Number of output neurons (default: 1)
W1: Weight matrix from input to hidden layer (inputSize x hiddenSize)
W2: Weight matrix from hidden to output layer (hiddenSize x outputSize)

"""

```
def __init__(self, input_size=3, hidden_size=4, output_size=1):  
    """Initialize the neural network.
```

Args:

input_size: Number of input features
hidden_size: Number of neurons in hidden layer
output_size: Number of output neurons

"""

```
super(FFNeuralNetwork, self).__init__()
```

```
self.inputSize = input_size  
self.hiddenSize = hidden_size  
self.outputSize = output_size
```

```
# Initialize weights with random values from normal distribution  
self.W1 = nn.Parameter(torch.randn(self.inputSize, self.hiddenSize))  
self.W2 = nn.Parameter(torch.randn(self.hiddenSize, self.outputSize))
```

```
# Intermediate variables for forward pass  
self.z = None # Input to hidden layer (before activation)  
self.z2 = None # Hidden layer output (after activation)  
self.z3 = None # Input to output layer (before activation)
```

```
# Intermediate variables for backward pass  
self.z_activation = None  
self.z_activation_derivative = None  
self.out_error = None  
self.out_delta = None  
self.z2_error = None  
self.z2_delta = None
```

```
def activation(self, z):
```

"""Apply sigmoid activation function.

Args:

z: Input tensor

```

>Returns:
    Tensor: Activated output
"""
self.z_activation = ActivationFunction.sigmoid(z)
return self.z_activation

def activation_derivative(self, z):
    """Compute derivative of sigmoid activation.

Args:
    z: Input tensor (typically activated output)

>Returns:
    Tensor: Derivative of activation
"""
self.z_activation_derivative = ActivationFunction.sigmoid_derivative(z)
return self.z_activation_derivative

def forward(self, X):
    """Perform forward propagation through the network.

Args:
    X: Input tensor of shape (batch_size, inputSize)

>Returns:
    Tensor: Output of the network, shape (batch_size, outputSize)
"""
# Layer 1: Input to Hidden
self.z = torch.matmul(X, self.W1) # z = X @ W1
self.z2 = self.activation(self.z) # z2 = sigmoid(z)

# Layer 2: Hidden to Output
self.z3 = torch.matmul(self.z2, self.W2) # z3 = z2 @ W2
output = self.activation(self.z3) # output = sigmoid(z3)

return output

def backward(self, X, y, output, learning_rate):
    """Perform backward propagation and update weights.

Args:
    X: Input tensor of shape (batch_size, inputSize)
    y: Target tensor of shape (batch_size, outputSize)
    output: Predicted output from forward pass
    learning_rate: Learning rate for weight updates
"""
# Output layer error and delta

```

```

        self.out_error = y - output # Error at output
        self.out_delta = self.out_error * self.activation_derivative(output)

        # Hidden layer error and delta
        self.z2_error = torch.matmul(self.out_delta, torch.t(self.W2))
        self.z2_delta = self.z2_error * self.activation_derivative(self.z2)

        # Update weights
        with torch.no_grad():
            self.W1 += torch.matmul(torch.t(X), self.z2_delta) * learning_rate
            self.W2 += torch.matmul(torch.t(self.z2), self.out_delta) * learning_rate
    ↵learning_rate

    def train(self, X, y, learning_rate):
        """Train the network for one epoch.

        Args:
            X: Input tensor of shape (batch_size, inputSize)
            y: Target tensor of shape (batch_size, outputSize)
            learning_rate: Learning rate for weight updates
        """
        # Forward pass
        output = self.forward(X)

        # Backward pass
        self.backward(X, y, output, learning_rate)

    @staticmethod
    def save_weights(model, path):
        """Save model state_dict to file.

        This method saves only the weights (state_dict) instead of the entire
        model object, which is more secure and avoids pickle-related issues.

        Args:
            model: FFNeuralNetwork instance to save
            path: File path to save the state_dict
        """
        torch.save(model.state_dict(), path)
        print(f"Model state_dict saved to {path}")

    @staticmethod
    def load_weights(path, input_size=3, hidden_size=4, output_size=1):
        """Load model weights from file.

        This method creates a new model instance and loads the saved state_dict.
        This approach avoids pickle security issues and is compatible with

```

PyTorch 2.6+ default security settings.

Args:

path: File path to load the state_dict from
input_size: Number of input features (must match saved model)
hidden_size: Number of neurons in hidden layer (must match saved_
model)
output_size: Number of output neurons (must match saved model)

Returns:

FFNeuralNetwork: Model instance with loaded weights

"""

Create a new model instance with the same architecture
model = FFNeuralNetwork(input_size=input_size,
 hidden_size=hidden_size,
 output_size=output_size)

Load the state_dict

state_dict = torch.load(path, weights_only=True)
model.load_state_dict(state_dict)

print(f"Model state_dict loaded from {path}")
return model

def predict(self, x_predict):

"""Make prediction on input data.

Args:

x_predict: Input tensor for prediction

Returns:

Tensor: Predicted output

"""

print("Predict data based on trained weights:")
print("Input:\n", x_predict)
output = self.forward(x_predict)
print("Output:\n", output)
return output

2.5 STEP 9 — Prepare Sample Data

Create and normalize training data (X , y) and test data ($x_predict$).

[40]: # Training data

```
X = torch.tensor(([2, 9, 0],  
                 [1, 5, 1],  
                 [3, 6, 2]), dtype=torch.float) # 3 x 3 tensor
```

```

y = torch.tensor([[90],
                 [100],
                 [88]], dtype=torch.float) # 3 x 1 tensor

print("Original X:")
print(X)
print("\nOriginal y:")
print(y)

# Normalize X by dividing each column by its max value
X_max, _ = torch.max(X, 0)
X = torch.div(X, X_max)
print("\nNormalized X (divided by column max):")
print(X)

# Normalize y by dividing by 100 (max test score)
y = y / 100
print("\nNormalized y (divided by 100):")
print(y)
print(X)

```

Original X:

```
tensor([[2., 9., 0.],
       [1., 5., 1.],
       [3., 6., 2.]])
```

Original y:

```
tensor([[ 90.],
       [100.],
       [ 88.]])
```

Normalized X (divided by column max):

```
tensor([[0.6667, 1.0000, 0.0000],
       [0.3333, 0.5556, 0.5000],
       [1.0000, 0.6667, 1.0000]])
```

Normalized y (divided by 100):

```
tensor([[0.9000],
       [1.0000],
       [0.8800]])
tensor([[0.6667, 1.0000, 0.0000],
       [0.3333, 0.5556, 0.5000],
       [1.0000, 0.6667, 1.0000]])
```

```
[41]: # Prediction data
x_predict = torch.tensor([[3, 8, 4]], dtype=torch.float) # 1 x 3 tensor
```

```

print("Original x_predict:")
print(x_predict)

# Normalize x_predict
x_predict = torch.div(x_predict, X_max)
print("\nNormalized x_predict:")
print(x_predict)
print(X_max)

```

```

Original x_predict:
tensor([[3., 8., 4.]])

Normalized x_predict:
tensor([[1.0000, 0.8889, 2.0000]])
tensor([3., 9., 2.])

```

2.6 STEP 10 — Training Loop

Train the network for 1000 epochs, print loss after each epoch.

```

[42]: # Create neural network instance
model = FFNeuralNetwork(input_size=3, hidden_size=4, output_size=1)

# Training parameters
num_epochs = 1000
learning_rate = 0.1

print("Starting training...")
print(f"Number of epochs: {num_epochs}")
print(f"Learning rate: {learning_rate}\n")

# Training loop
for epoch in range(num_epochs):
    # Forward pass to compute loss
    output = model(X)
    loss = torch.mean((y - output) ** 2).detach().item()

    # Print loss every 100 epochs
    if (epoch + 1) % 100 == 0 or epoch == 0:
        print(f"Epoch #{epoch + 1} - Loss: {loss:.6f}")

    # Training step
    model.train(X, y, learning_rate)

print("\nTraining completed!")

```

```

Starting training...
Number of epochs: 1000
Learning rate: 0.1

```

```
Epoch #1 - Loss: 0.374021
Epoch #100 - Loss: 0.030377
Epoch #200 - Loss: 0.009770
Epoch #300 - Loss: 0.005774
Epoch #400 - Loss: 0.004424
Epoch #500 - Loss: 0.003831
Epoch #600 - Loss: 0.003528
Epoch #700 - Loss: 0.003356
Epoch #800 - Loss: 0.003249
Epoch #900 - Loss: 0.003179
Epoch #1000 - Loss: 0.003129
```

Training completed!

```
[43]: model.predict(x_predict)
```

```
Predict data based on trained weights:
Input:
tensor([[1.0000, 0.8889, 2.0000]])
Output:
tensor([[0.9303]], grad_fn=<MulBackward0>)
```

```
[43]: tensor([[0.9303]], grad_fn=<MulBackward0>)
```

```
[44]: p = model.predict(x_predict)
```

```
print(f"\nPredicted value (scaled): {p.item():.6f}")
```

```
Predict data based on trained weights:
Input:
tensor([[1.0000, 0.8889, 2.0000]])
Output:
tensor([[0.9303]], grad_fn=<MulBackward0>)
```

Predicted value (scaled): 0.930308

```
[45]: model.state_dict()
```

```
[45]: OrderedDict([('W1',
                    tensor([[ 1.0158e+00,  2.8769e-01, -4.5191e-01, -1.5456e+00],
                           [ 1.4163e-03,  4.6565e-01, -6.5942e-01, -2.3856e-01],
                           [ 1.5531e+00, -5.8867e-01, -8.9156e-01,  4.2244e-01]]),
                   ('W2',
                    tensor([[1.9637],
                           [1.2179],
                           [0.2600],
                           [0.6201]])))])
```

2.7 STEP 8 — Save Model

Save the trained model's state_dict (weights only) to file. This approach is more secure and compatible with PyTorch 2.6+ security defaults, avoiding pickle-related issues.

```
[46]: print(model.state_dict().keys())
odict_keys(['W1', 'W2'])
```

```
[47]: # Save trained model
model_path = "ffnn_model.pth"
FFNeuralNetwork.save_weights(model, model_path)
```

Model state_dict saved to ffnn_model.pth

2.8 STEP 11 — Prediction

Load saved state_dict by creating a new model instance with matching architecture, then load the weights. Perform prediction on new data.

```
[48]: # Load saved model
# Note: Must specify architecture parameters matching the saved model
loaded_model = FFNeuralNetwork.load_weights(model_path,
                                              input_size=3,
                                              hidden_size=4,
                                              output_size=1)

# Make prediction
prediction = loaded_model.predict(x_predict)
print(f"\nPredicted value (scaled): {prediction.item():.6f}")
print(f"Predicted value (original scale): {prediction.item() * 100:.2f}")
```

Model state_dict loaded from ffnn_model.pth

Predict data based on trained weights:

Input:
tensor([[1.0000, 0.8889, 2.0000]])
Output:
tensor([[0.9303]], grad_fn=<MulBackward0>)

Predicted value (scaled): 0.930308

Predicted value (original scale): 93.03

2.9 STEP 12 — Experiments

Experiment with different hyperparameters to observe changes in loss and output.

```
[49]: # Experiment 1: Different learning rates
print("=" * 60)
print("Experiment 1: Testing different learning rates")
print("=" * 60)
```

```

learning_rates = [0.01, 0.1, 0.5, 1.0]
for lr in learning_rates:
    model_exp = FFNeuralNetwork(input_size=3, hidden_size=4, output_size=1)

    # Train for 500 epochs
    for epoch in range(500):
        model_exp.train(X, y, lr)

    output = model_exp(X)
    loss = torch.mean((y - output) ** 2).detach().item()
    print(f"Learning rate: {lr:.4f} - Final Loss: {loss:.6f}")

print()

```

```

=====
Experiment 1: Testing different learning rates
=====

Learning rate: 0.01 - Final Loss: 0.063629
Learning rate: 0.10 - Final Loss: 0.004101
Learning rate: 0.50 - Final Loss: 0.002876
Learning rate: 1.00 - Final Loss: 0.003301

```

2.9.1 Config Hidden layer

```

[50]: # Experiment 2: Different number of hidden neurons
print("=" * 60)
print("Experiment 2: Testing different hidden layer sizes")
print("=" * 60)

hidden_sizes = [2, 4, 6, 8, 16]
for hidden_size in hidden_sizes:
    model_exp = FFNeuralNetwork(input_size=3, hidden_size=hidden_size, ↴
                                output_size=1)

    # Train for 1000 epochs
    for epoch in range(1000):
        model_exp.train(X, y, 0.1)

    output = model_exp(X)
    loss = torch.mean((y - output) ** 2).detach().item()
    print(f"Hidden size: {hidden_size:2d} - Final Loss: {loss:.6f}")

print()

```

```

=====
Experiment 2: Testing different hidden layer sizes
=====
```

```
=====
Hidden size: 2 - Final Loss: 0.003283
Hidden size: 4 - Final Loss: 0.003004
Hidden size: 6 - Final Loss: 0.004316
Hidden size: 8 - Final Loss: 0.003604
Hidden size: 16 - Final Loss: 0.003059
```

```
[51]: # Experiment 3: Compare predictions with different configurations
print("=" * 60)
print("Experiment 3: Predictions with different models")
print("=" * 60)

configs = [
    {"hidden_size": 4, "lr": 0.1, "name": "Default (4 neurons, lr=0.1)" },
    {"hidden_size": 8, "lr": 0.1, "name": "More neurons (8 neurons, lr=0.1)" },
    {"hidden_size": 4, "lr": 0.5, "name": "Higher LR (4 neurons, lr=0.5)" },
]

for config in configs:
    model_exp = FFNeuralNetwork(
        input_size=3,
        hidden_size=config["hidden_size"],
        output_size=1
    )

    # Train for 1000 epochs
    for epoch in range(1000):
        model_exp.train(X, y, config["lr"])

    prediction = model_exp.predict(x_predict)
    print(f'{config["name"]}: {prediction.item() * 100:.2f}')
    print()
```

```
=====
Experiment 3: Predictions with different models
=====
```

Predict data based on trained weights:

Input:
`tensor([[1.0000, 0.8889, 2.0000]])`

Output:
`tensor([[0.9763]], grad_fn=<MulBackward0>)`

Default (4 neurons, lr=0.1): 97.63

Predict data based on trained weights:

Input:
`tensor([[1.0000, 0.8889, 2.0000]])`

Output:

```
tensor([[0.9748]], grad_fn=<MulBackward0>)
More neurons (8 neurons, lr=0.1): 97.48
```

Predict data based on trained weights:

Input:
tensor([[1.0000, 0.8889, 2.0000]])

Output:
tensor([[0.9453]], grad_fn=<MulBackward0>)
Higher LR (4 neurons, lr=0.5): 94.53